

(COM4523-B) **ARTIFICIAL** **NEURAL** **NETWORKS**

CODE EXPLANATIONS

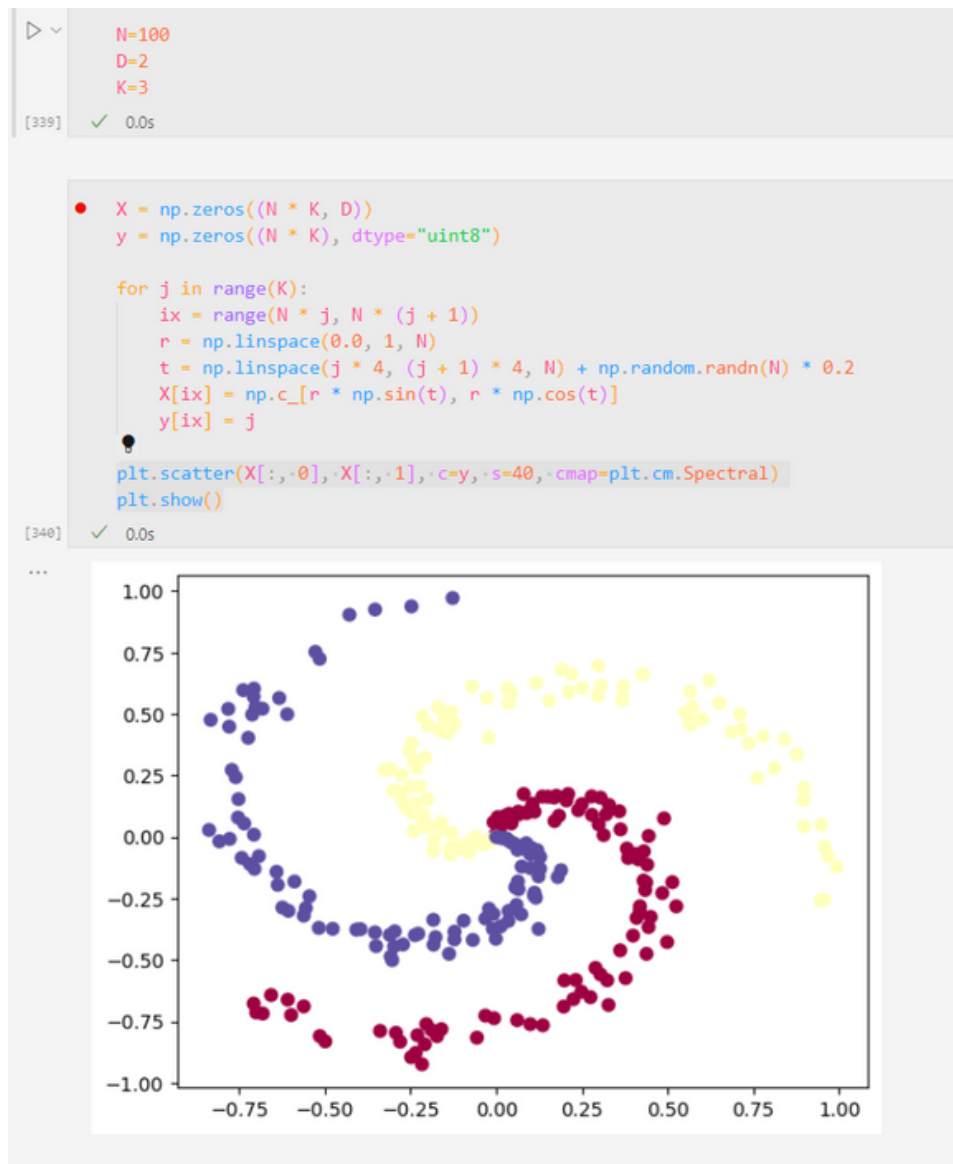
Prepared By

Selim Koç

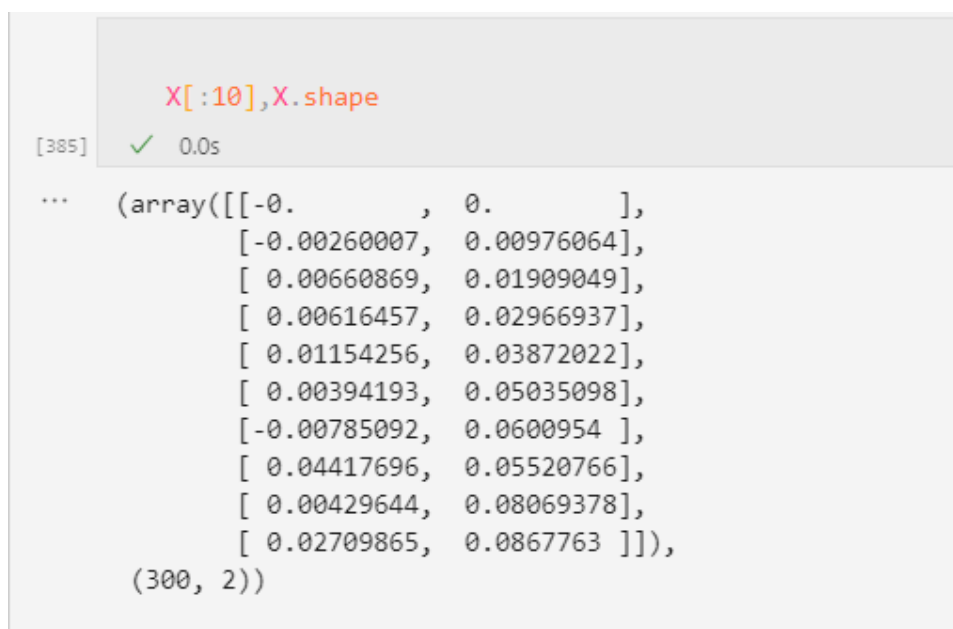
19290317

Softmax Linear Classifier

Firstly we create our dataset for the model



Let's examine the content and shape of our data.



[illegible]

Define a function to separate the data into train ,validate and test sets

```
def random_split_data(X, y, train_percentage=0.7, val_percentage=0.15, test_percentage=0.15):
    """
    Randomly splits the data into training, validation, and test sets.

    Parameters:
    - X: Feature matrix
    - y: Labels
    - train_percentage: Percentage of data for training (default is 0.7)
    - val_percentage: Percentage of data for validation (default is 0.15)
    - test_percentage: Percentage of data for testing (default is 0.15)

    Returns:
    - X_train, y_train: Training set
    - X_val, y_val: Validation set
    - X_test, y_test: Test set
    """

    # Check if percentages sum up to 1
    total_percentage = train_percentage + val_percentage + test_percentage
    if total_percentage != 1.0:
        raise ValueError("Percentages should sum up to 1.0")

    # Get the total number of examples
    total_examples = len(X)

    # Shuffle the data
    indices = np.arange(total_examples)
    np.random.shuffle(indices)

    # Split data into training, validation, and test sets
    train_size = int(train_percentage * total_examples)
    val_size = int(val_percentage * total_examples)

    train_indices = indices[:train_size]
    remaining_indices = indices[train_size:]

    val_indices = remaining_indices[:val_size]
    test_indices = remaining_indices[val_size:]

    # Create sets
    X_train, y_train = X[train_indices], y[train_indices]
    X_val, y_val = X[val_indices], y[val_indices]
    X_test, y_test = X[test_indices], y[test_indices]

    return X_train, y_train, X_val, y_val, X_test, y_test
```


Separate the data using “random_split_data” function

```
[342] ✓ 0.0s
X_train, y_train, X_val, y_val, X_test, y_test = random_split_data(X, y)
```

Initialize the weight matrix W_{exp} with random values from a standard normal distribution

```
[343] ✓ 0.0s
W_exp=np.random.randn(D, K)
W_exp #W_exp matrix has dimensions D (number of features) by K (number of classes)
... array([[ 0.19054612, -0.65011432, -0.12252903],
           [ 0.66658672,  0.53478674, -0.42044854]])
```

We use a small scaling factor (0.01) to initialize the weights (W) in order to prevent issues like the gradients becoming too high during training

```
[344] ✓ 0.0s
W = 0.01 * W_exp
W
... array([[ 0.00190546, -0.00650114, -0.00122529],
           [ 0.00666587,  0.00534787, -0.00420449]])
```

Initialize the bias vector b with zeros.

```
[345] ✓ 0.0s
b = np.zeros((1, K))
b # The vector has dimensions 1 by K, where K is the number of classes.
... array([[0., 0., 0.]])
```

Calculate the scores for each class by performing a dot product between the input features (X) and the weight matrix (W),and then adding the bias vector (b).

```
[346] ✓ 0.0s
scores = np.dot(X_train, W) + b
scores[:10],scores.shape
... (array([[ -4.80355658e-06, -7.04177790e-04,  9.86941533e-07],
           [ 1.86070408e-03,  3.53705407e-03, -1.16767298e-03],
           [-6.10878849e-04, -3.97662089e-03,  3.75140443e-04],
           [-8.33088978e-04, -1.65696371e-03,  5.22585779e-04],
           [ 1.86949002e-04,  2.10739073e-03, -1.12207903e-04],
           [-2.39780453e-03, -2.02744055e-03,  1.51210888e-03],
           [ 3.04645074e-03, -3.66833765e-03, -1.93937448e-03],
           [-2.83402249e-04,  3.08273823e-04,  1.80318032e-04],
           [ 1.35025703e-03,  4.58403652e-06, -8.54819164e-04],
           [-1.40804762e-03,  5.59862021e-03,  9.07750672e-04]]),
      (210, 3))
```

We create a variable to use in loops later in the code which contains the number of rows of our input data

```
num_examples = X_train.shape[0] # This variable is created to be used in loops later in the code.
num_examples

[347] ✓ 0.0s
... 210
```

Calculate the scores corresponding to the correct classes for each example.

```
correct_class_scores = scores[range(num_examples), y_train]
correct_class_scores[:10], correct_class_scores[-10:]

[348] ✓ 0.0s
... (array([ 9.86941533e-07,  3.53705407e-03, -6.10878849e-04,  5.22585779e-04,
            2.10739073e-03,  1.51210888e-03, -3.66833765e-03,  3.08273823e-04,
            1.35025703e-03,  9.07750672e-04]),
      array([ 0.00102005, -0.00131674, -0.00710276,  0.00319932,  0.00197148,
            0.00325313, -0.00559717, -0.00669607,  0.00217245,  0.00258924]))
```

Calculate the exponential of the scores to obtain unnormalized probabilities.

```
exp_scores = np.exp(scores)
exp_scores[:10], exp_scores[-10:]

[349] ✓ 0.0s
... (array([[0.99999952, 0.99929607, 1.00000099],
            [1.00186244, 1.00354332, 0.99883301],
            [0.99938931, 0.99603128, 1.00037521],
            [0.99916726, 0.99834441, 1.00052272],
            [1.00018697, 1.00210961, 0.9998878 ],
            [0.99760507, 0.99797461, 1.00151325],
            [1.0030511 , 0.99633838, 0.9980625 ],
            [0.99971664, 1.00030832, 1.00018033],
            [1.00135117, 1.00000458, 0.99914555],
            [0.99859294, 1.00561432, 1.00090816]]),
      array([[1.00102057, 0.99997735, 0.99935436],
            [1.00488224, 0.99868413, 0.99691755],
            [1.00108201, 0.99292241, 0.99929489],
            [1.00372097, 1.00320444, 0.99766074],
            [0.99690236, 1.00252489, 1.00197343],
            [1.00282522, 1.00325842, 0.99822498],
            [0.99441846, 0.99702854, 1.00354108],
            [1.00132942, 0.99332629, 0.99913976],
            [0.99657176, 0.99943059, 1.00217481],
            [1.00047676, 1.0025926 , 0.99970584]]))
```

Calculate the normalized probabilities by dividing the exponential scores by the sum of exponential scores for each example.

```
probs = exp_scores / np.sum(exp_scores, axis=1, keepdims=True)
probs[:10]

[350] ✓ 0.0s
... array([[0.33341039, 0.33317729, 0.33341232],
            [0.33348296, 0.33404246, 0.33247458],
            [0.33359727, 0.33247636, 0.33392637],
            [0.33327412, 0.33299965, 0.33372623],
            [0.33315308, 0.33379349, 0.33305343],
            [0.33285757, 0.33298087, 0.33416156],
            [0.33463458, 0.33239511, 0.33297031],
            [0.33321608, 0.33341329, 0.33337063],
            [0.33372796, 0.33327917, 0.33299287],
            [0.3322977 , 0.33463417, 0.33306813]])
```

Calculate the sum of probabilities for each example, ensuring that the total probability for each example sums to 1

[illegible]

Calculate the cross-entropy loss by summing the negative logarithm of the predicted probabilities for the correct classes.

```
loss = -np.sum(np.log(probs[range(num_examples), y_train])) / num_examples
loss
[352] ✓ 0.0s
... 1.098683179274858
```

This demonstrates that even if all values are randomly assigned, the worst-case loss can be determined by \log_2 of the number of classes.

```
import math
result = math.log2(3)
result
```

[353] ✓ 0.0s

... 1.584962500721156

Create a copy of the probabilities array to represent the gradient of the scores with respect to the loss.

```

dscores = probs.copy()
dscores[range(num_examples), y_train] -= 1 #Subtracting 1 from the probabilities of correct classes into
dscores

[354] ✓ 0.0s

... array([[ 0.33341039,  0.33317729, -0.66658768],
 [ 0.33348296, -0.66595754,  0.33247458],
 [-0.66640273,  0.33247636,  0.33392637],
 [ 0.33327412,  0.33299965, -0.66627377],
 [ 0.33315308, -0.66620651,  0.33305343],
 [ 0.33285757,  0.33298087, -0.66583844],
 [ 0.33463458, -0.66760489,  0.33297031],
 [ 0.33321608, -0.66658671,  0.33337063],
 [-0.66627204,  0.33327917,  0.33299287],
 [ 0.3322977 ,  0.33463417, -0.66693187],
 [ 0.33417994, -0.66619912,  0.33201917],
 [ 0.33421139, -0.66835249,  0.3341411 ],
 [ 0.33465896, -0.66648381,  0.33182485],
 [-0.66609129,  0.33266172,  0.33342957],
 [ 0.3329997 ,  0.33293455, -0.66593424],
 [-0.66646164,  0.33331936,  0.33314227],
 [-0.66618765,  0.33299132,  0.33319633],
```

Calculate the gradient of the weights dW :

Multiply these probabilities (dscores) with the transpose of our input data x , we would essentially be applying a weighted sum to each row of x based on the calculated probabilities for each class.

Each row corresponds to a class, and the columns represent the weighted sum of the input features based on the calculated probabilities for that class.

Calculate the gradient of the bias terms db :

Summing the probabilities (dscores) along the rows ($axis=0$) provides the contribution of each class to the gradient of the bias terms. Each element in the resulting array corresponds to the sum of probabilities for the respective class across all input samples.

```
[355] ✓ 0.0s  
dW = np.dot(X_train.T, dscores)  
db = np.sum(dscores, axis=0, keepdims=True)
```

```
[356] ✓ 0.0s  
dW  
... array([[ -9.09618303, -10.36651998,  19.46270301],  
          [ 16.18079783, -19.53639229,   3.35559446]])
```

```
[357] ✓ 0.0s  
db  
... array([[ 1.98814908,  1.0052981 , -2.99344718]])
```

Update the weights (W) and biases (b) using the learning rate (α)

```
[358] ✓ 0.0s  
alpha = 0.01  
W -= alpha * dW  
b -= alpha * db
```

New weights and biases

```
[359] ✓ 0.0s  
W,b  
... (array([[ 0.09286729,  0.09716406, -0.19585232],  
          [-0.15514211,  0.20071179, -0.03776043]]),  
     array([[ -0.01988149, -0.01005298,  0.02993447]]))
```

We create a function for training our model

```
def trainModel(X_train, y_train, X_val, y_val, W, b, epochs=400, alpha=0.01):
    """
    Train a model using gradient descent.

    Parameters:
    - X_train: Training feature matrix
    - y_train: Training true labels
    - X_val: Validation feature matrix
    - y_val: Validation true labels
    - W: Model weights
    - b: Model bias
    - epochs: Number of training epochs (default is 1000)
    - alpha: Learning rate (default is 0.01)

    Returns:
    - W: Updated weights after training
    - b: Updated bias after training
    - train_losses: List of training losses for each epoch
    - val_losses: List of validation losses for each epoch
    """
    # Lists to store losses for plotting. We can use these to make sure we don't overfit the training data.
    train_losses = []
    val_losses = []

    # Training Loop
    for epoch in range(epochs):
        # Forward pass on the training set
        scores_train = np.dot(X_train, W) + b
        exp_scores_train = np.exp(scores_train)
        probs_train = exp_scores_train / np.sum(exp_scores_train, axis=1, keepdims=True)
        loss_train = -np.sum(np.log(probs_train[range(len(y_train)), y_train])) / len(y_train)

        # Backward pass on the training set
        dscores_train = probs_train.copy()
        dscores_train[range(len(y_train)), y_train] -= 1
        dW_train = np.dot(X_train.T, dscores_train)
        db_train = np.sum(dscores_train, axis=0, keepdims=True)

        # Update parameters for the training set
        W -= alpha * dW_train
        b -= alpha * db_train

        # Store the training loss for plotting
        train_losses.append(loss_train)

        # Forward pass on the validation set
        scores_val = np.dot(X_val, W) + b
        exp_scores_val = np.exp(scores_val)
        probs_val = exp_scores_val / np.sum(exp_scores_val, axis=1, keepdims=True)
        loss_val = -np.sum(np.log(probs_val[range(len(y_val)), y_val])) / len(y_val)

        # Store the validation loss for plotting
        val_losses.append(loss_val)

    return W, b, train_losses, val_losses
```

Update the weights (W) and biases (b) using the function . We use both trained and validation data because we try to see our model gonna be overfit or underfit . We change the epoch value 1000 to 400 because we easily see there is no change on the line on loss function and also our data is so small for 1000.

```
Wtrained, btrained, train_losses, val_losses = trainModel(X_train, y_train, X_val, y_val, W, b, epochs=1000, alpha=0.01)
train_losses[:10], val_losses[:10]
```

[illegible]

Checking accuracy using the model with the test set and updated weights and biases.

```
def AccuracyCheck(X, y, W, b):
    """
    Evaluate the accuracy of the model on a given dataset.

    Parameters:
    - X: Feature matrix
    - y: True labels
    - W: Weights
    - b: Bias

    Returns:
    - accuracy: Accuracy of the model on the given dataset
    """
    scores = np.dot(X, W) + b
    # Determine predicted classes (class with the highest probability)
    predicted_classes = np.argmax(scores, axis=1)
    # Count the number of correct predictions
    correct_prediction_count = np.sum(predicted_classes == y)
    # Calculate the overall accuracy rate
    accuracy = correct_prediction_count / X.shape[0]
    return accuracy

[362] ✓ 0.0s

accuracyOnTest = AccuracyCheck(X_test, y_test, WTrained, bTrained)
accuracyOnTest

[370] ✓ 0.0s

... 0.6
```

We also made a functions for showing decision boundaries and our datas

```
def plot_decision_boundary(X, y, W, b):
    """
    Plot the decision boundary of a trained model along with the data points.

    Parameters:
    - X: Feature matrix
    - y: True labels
    - W: Model weights
    - b: Model bias
    """
    # Create a meshgrid of points
    h = 0.02 # step size in the mesh
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))

    # Predict the class labels for each point in the meshgrid
    Z = np.dot(np.c_[xx.ravel(), yy.ravel()], W) + b
    Z = np.argmax(Z, axis=1)
    Z = Z.reshape(xx.shape)

    # Plot the decision boundary
    plt.contourf(xx, yy, Z, cmap=plt.cm.Spectral, alpha=0.8)
    plt.scatter(X[:, 0], X[:, 1], c=y, s=40, cmap=plt.cm.Spectral, edgecolors='k')
    plt.xlabel('Feature 1')
    plt.ylabel('Feature 2')
    plt.title('Decision Boundary')

[364] ✓ 0.0s

def plot_data(features, output, type):
    plt.scatter(features[:, 0], features[:, 1], c=output, s=40, cmap=plt.cm.Spectral, edgecolors='k')
    plt.title(f'{type} Data')
    plt.show()

[390] ✓ 0.0s
```


Then we plot the decision boundary, and our data sets

▶

```
# Plot decision boundary on training set  
plot_decision_boundary(X_train, y_train, WTrained, bTrained)  
plt.show()
```

[418] ✓ 0.1s


...





```
plot_data(X_train, y_train, "Train")  
plot_data(X_val, y_val, "Validation")  
plot_data(X_test, y_test, "Test")
```

[419] ✓ 0.2s

...



...




Finally we can see our model's loss

```
plt.plot(train_losses, label='Training Loss')
plt.plot(val_losses, label='Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training and Validation Loss Curves')
plt.legend()
plt.show()
```

[481] ✓ 0.1s

...



Neural Network

Firstly we create our data and divide into train test and validation set using `random_split_data` function which we mentioned before

```
154
155 # Generating synthetic data
156 N = 100
157 D = 2
158 K = 3
159 X = np.zeros((N * K, D))
160 y = np.zeros((N * K), dtype="uint8")
161
162 for j in range(K):
163     ix = range(N * j, N * (j + 1))
164     r = np.linspace(0.0, 1, N)
165     t = np.linspace(j * 4, (j + 1) * 4, N) + np.random.randn(N) * 0.2
166     X[ix] = np.c_[r * np.sin(t), r * np.cos(t)]
167     y[ix] = j
168
169 plt.scatter(X[:, 0], X[:, 1], c=y, s=40, cmap=plt.cm.Spectral)
170 plt.show()
171
172 # Splitting the data into training, validation, and test sets
173 X_train, y_train, X_val, y_val, X_test, y_test = random_split_data(X, y)
```

Since we will create a 2-layer neural network we create 2 weight and bias

```
35 # Function to initialize neural network parameters
36 def initialize_parameters(D, neuron_count, K):
37     # Initialize weights W1 with small random values.
38     W1 = 0.01 * np.random.randn(D, neuron_count)
39
40     # Initialize biases b1 with zeros.
41     b1 = np.zeros((1, neuron_count))
42
43     # Initialize weights W2 with small random values.
44     W2 = 0.01 * np.random.randn(neuron_count, K)
45
46     # Initialize biases b2 with zeros.
47     b2 = np.zeros((1, K))
48
49     return W1, b1, W2, b2
```

```
174
175 # Initializing parameters for the neural network
176 W1, b1, W2, b2 = initialize_parameters(D, 100, K)
177
```

In `forward_pass` firstly, we calculate the dot product of input and weight matrix then add the bias to find output then we use the `relu` activation function for feature transform and we get the hidden layer value. Then we apply same calculation on hidden layer using `W2` and `b2` we get a final output. Lastly we apply `Softmax` to get probabilities

```
51 # Forward Pass Function
52 def forward_pass(X, W1, b1, W2, b2):
53     # Calculate the first layer output z1 and apply ReLU activation function.
54     output1 = np.dot(X, W1) + b1
55     hidden = np.maximum(0, output1) # ReLU activation
56
57     # Calculate the second layer output z2.
58     finalOutput = np.dot(hidden, W2) + b2
59
60     # Apply the softmax function to get probabilities.
61     exp_scores = np.exp(finalOutput)
62     probs = exp_scores / np.sum(exp_scores, axis=1, keepdims=True)
63
64     return hidden, probs
65
```

We call the loss function with the output of forward pass which is probabilities and calculate the cross-entropy loss in the calculate_loss function

```
66 # Loss Calculation Function
67 def calculate_loss(probs, y):
68     # Calculate cross-entropy loss for each class and sum them for the total loss.
69     num_examples = len(y)
70     correct_logprobs = -np.log(probs[range(num_examples), y])
71     loss = np.sum(correct_logprobs) / num_examples
72
73     return loss
```

```
178 # Forward pass to calculate loss before training
179 hidden, probs = forward_pass(X_train, W1, b1, W2, b2)
180 loss = calculate_loss(probs, y_train)
181 print("Loss on the first iteration: ", loss)
182
```

In the backward pass function, dscores is initially created as a copy of the probabilities (probs) to ensure that the probabilities remain unchanged throughout the subsequent computations. To expedite learning, the loss is increased by directly modifying dscores for the correct class, achieved by subtracting 1 from the probability associated with the correct class for each example. This manipulation effectively boosts the gradient, potentially accelerating the learning process.

The gradients for the second layer parameters, dW2 and db2, are then computed by utilizing the transposed hidden layer (hidden.T) and the modified dscores. Subsequently, the gradient of the loss with respect to the output of the first layer, da1, is calculated. This involves multiplying dscores by the transposed weights of the second layer (W2.T), with an additional application of the ReLU activation function gradient by setting values to zero where the corresponding values in the hidden layer are less than or equal to 0.

The gradients for the first layer parameters, dW1 and db1, are determined using the transposed input data (X.T) and the previously calculated da1. Following this, both layers' parameters are updated using gradient descent: the learning rate is multiplied by the corresponding gradients, and the results are subtracted from the current parameter values.

```
74
75 # Backward Pass and Parameter Update Function
76 def backward_pass(X, y, hidden, probs, W1, W2, b1, b2, learning_rate):
77     # Compute gradients.
78     num_examples = len(y)
79     dscores = probs.copy()
80     dscores[range(num_examples), y] -= 1
81
82     # Compute gradients for the second layer parameters.
83     dW2 = np.dot(hidden.T, dscores)
84     db2 = np.sum(dscores, axis=0, keepdims=True)
85
86     # Compute gradients for the first layer output.
87     da1 = np.dot(dscores, W2.T)
88     da1[hidden <= 0] = 0 # ReLU activation gradient
89
90     # Compute gradients for the first layer parameters.
91     dW1 = np.dot(X.T, da1)
92     db1 = np.sum(da1, axis=0, keepdims=True)
93
94     # Update parameters.
95     W1 -= learning_rate * dW1
96     b1 -= learning_rate * db1
97     W2 -= learning_rate * dW2
98     b2 -= learning_rate * db2
99
100     return W1, b1, W2, b2
101
```


After updating our values, we find our new loss by repeating the previous process and we see that we have a lower loss.

```
104
183 # Backward Pass and Parameter Update
184 W1, b1, W2, b2 = backward_pass(X_train, y_train, hidden, probs, W1, W2, b1, b2, learning_rate=0.01)
185
186 # After the first iteration, calculate loss again to see if it decreases
187 hidden, probs = forward_pass(X_train, W1, b1, W2, b2)
188 loss = calculate_loss(probs, y_train)
189 print("Loss after one backpropagation: ", loss)
190
```

Now we come to the most important part we define a function which is `train_neural_network`.

This function takes in both a training set (`X_train, y_train`) and a validation set (`X_val, y_val`), along with the initial parameters of the neural network (`W1, b1, W2, b2`), the learning rate (`learning_rate`), and the number of training epochs (`epochs`). Within the function, there's a loop that iterates over the specified number of epochs. During each epoch:

Forward Pass: The function computes the forward pass on the training and validation set, generating the hidden layer activations (`hidden`) and output probabilities (`probs`).

Loss Calculation: It calculates the cross-entropy loss for the training and validation set using the `calculate_loss` function.

Backward Pass and Parameter Update: The backward pass is performed to compute gradients, and the neural network parameters (`W1, b1, W2, b2`) are updated using the `backward_pass` function.


At the end of each epoch, the training and validation losses are stored in separate lists (`train_losses` and `val_losses`), allowing for later analysis, such as plotting learning curves. The function eventually returns the final parameters of the neural network (`W1, b1, W2, b2`) after the training process, along with the lists of training and validation losses. This systematic approach enables monitoring the model's performance on both the training and validation sets throughout the training process.

```
102 # Training Neural Network Function
103 def train_neural_network(X_train, y_train, X_val, y_val, W1, b1, W2, b2, learning_rate, epochs):
104     # Store Losses for plotting.
105     train_losses = []
106     val_losses = []
107
108     for epoch in range(epochs):
109         # Forward pass for training set.
110         hidden, probs = forward_pass(X_train, W1, b1, W2, b2)
111
112         # Loss calculation for training set.
113         train_loss = calculate_loss(probs, y_train)
114         train_losses.append(train_loss)
115
116         # Backward pass and parameter update for training set.
117         W1, b1, W2, b2 = backward_pass(X_train, y_train, hidden, probs, W1, W2, b1, b2, learning_rate)
118
119         # Forward pass for validation set.
120         val_hidden, val_probs = forward_pass(X_val, W1, b1, W2, b2)
121
122         # Loss calculation for validation set.
123         val_loss = calculate_loss(val_probs, y_val)
124         val_losses.append(val_loss)
125
126     return W1, b1, W2, b2, train_losses, val_losses
127
```

```
190
191 # Training the neural network
192 W1, b1, W2, b2, train_losses, val_losses = train_neural_network(X_train, y_train, X_val, y_val, W1, b1, W2, b2, learning_rate=0.01, epochs=850)
193
```

Finally we plot all the necessary things

```
193
194 # Plot the training Loss curve
195 plt.plot(train_losses, label='Training Loss')
196 plt.plot(val_losses, label='Validation Loss')
197 plt.xlabel('Epoch')
198 plt.ylabel('Loss')
199 plt.title('Training and Validation Loss Curve')
200 plt.legend()
201 plt.show()
202
203
204
205 # Calculate accuracy on the test set
206 testAccuracy = calculate_accuracy(X_test, y_test, W1, b1, W2, b2)
207 print("Accuracy on the test set: {:.4%}".format(testAccuracy))
208
209
210 # Plot the learned decision boundaries on the training data
211 plot_decision_boundaries(X_train, y_train, W1, b1, W2, b2)
212
```



Loss on the first iteration: 1.0985255770082003
Loss after one backpropagation: 1.097886958650164
Accuracy on the test set: 97.7778%

